

T-SQL: Transactions and Concurrency Control

—

September, 2019

1 Transactions

2 Concurrency Controls

Section 1

Transactions

Atomicity

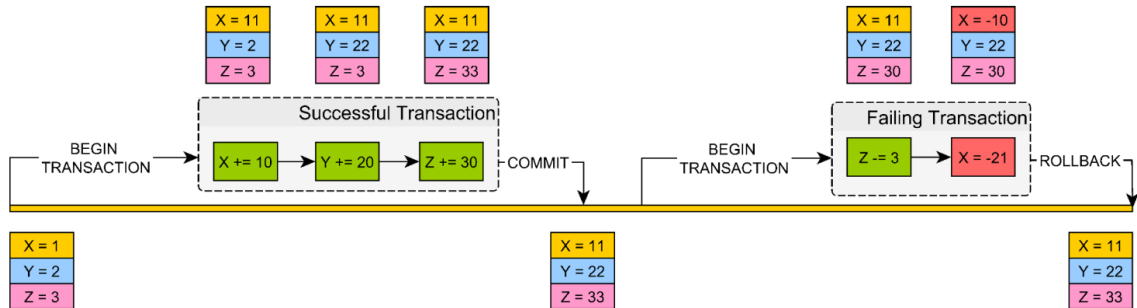


Figure 1: Atomic units of work

What is a Transaction (Xact)

- Collection of operations that form a single logical unit
- Logical unit:
 - begin transaction ... (SQL) end transaction
- Operations:
 - Read, Write
 - Special actions: begin, commit, rollback

ACID properties

- Atomicity
 - All action in the Xact happen, or none happen
- Consistency
 - Consistent DB + consistent Xact \Rightarrow consistent DB
- Isolation
 - Execution of one Xact is isolated from that of other Xacts
 - Concurrency control
- Durability
 - If a Xact commits, its effects persist
 - Recovery management

Autocommit mode in SQL Server (default)

- If not explicitly start a transaction
 - Each statement treated as a separated tran.
 - If error, automatically rollbacked.
 - Otherwise, automatically committed.

When to use explicit transactions?

- Two or more action queries affecting related data
- Update foreign key references
- Move rows from one table to another
- Action query whose values inserted depend on previously run select query
- A failure of any set of sql statements would violate integrity

Example: a transaction to insert invoices and items

```
declare @InvoiceID int;
begin try
    begin tran;
        insert Invoices
            values (34, 'ZXA-080', '2016-04-30', 14092.59,
                0, 0, 3, '2016-05-30', NULL);
        set @InvoiceID = @@IDENTITY;
        insert InvoiceLineItems values (@InvoiceID, 1, 160, 4447.23,
            'HW upgrade');
        insert InvoiceLineItems values (@InvoiceID, 2, 167, 9645.36,
            'OS upgrade');
    commit tran;
    print 'transaction committed';
end try
```

Example: a transaction to insert invoices and items (cont.)

```
begin catch
    print 'error when inserting, rolling back transaction';
    rollback tran;
end catch;
go
```

T-SQL statement for processing transactions

```
begin {tran | transactions}  
save {tran | transactions} 'save_point_name'  
commit [tran | transactions]  
rollback [[tran | transactions] ['save_point_name']]
```

Nested transactions

- No true nested transactions
- Use @@trancount counter

BEGIN TRAN

- @@trancount += 1

COMMIT TRAN

- @@trancount == 1: commit all trans, @@trancount = 0
- @@trancount > 1: nothing committed, @@trancount -= 1

ROLLBACK TRAN

- rollback all trans regardless of nesting level
- @@trancount = 0

Example: nested transactions

```
begin tran;
    print 'First Tran @@trancount: ' + convert(varchar,@@trancount);
    delete Invoices;
    begin tran;
        print 'Second Tran @@trancount:' + convert(varchar,@@trancount);
        delete Vendors;
    commit tran;  -- this commit decrements @@trancount
                  -- it doesn't commit 'delete Vendors'
    print 'commit @@trancount: ' + convert(varchar,@@trancount);
rollback tran;
```

Save points

- Create

```
SAVE TRAN 'save_point_name'
```

- Rollback upto & including the save_point

```
ROLLBACK TRAN 'save_point_name'
```

- Rollback entire transaction

```
ROLLBACK TRAN
```

Example: a transaction with save_points

```
begin tran
    delete #VendorCopy where VendorID = 1;
    save tran Vendor1;  --1st save_point
        delete #VendorCopy where VendorID = 2;
        save tran Vendor2; --2nd save_point
            delete #VendorCopy where VendorID = 3;
            select * from #VendorCopy;
        rollback tran Vendor2;  --rollback 2nd save_point
        select * from #VendorCopy;
    rollback tran Vendor1; --rollback 1st save_point
    select * from #VendorCopy;
commit tran
select * from #VendorCopy;
go
```

Section 2

Concurrency Controls

Problems

- Dirty reads
- Lost updates
- Nonrepeatable reads
- Phantom reads

Dirty reads

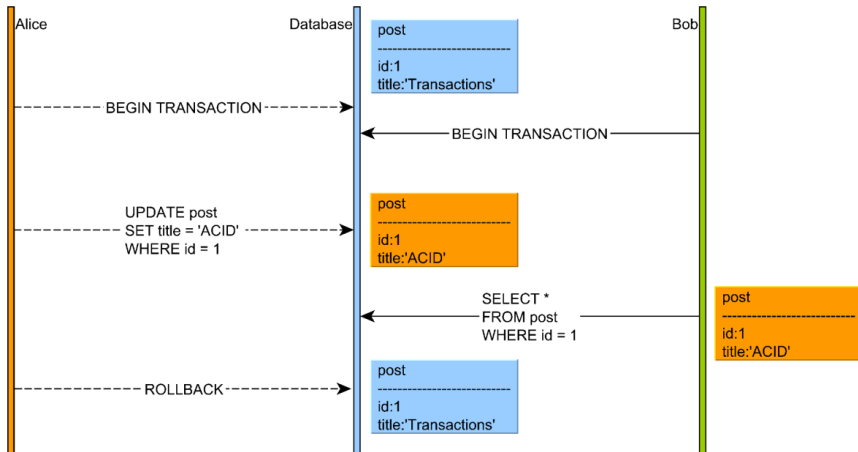


Figure 2: Dirty read

Lost updates

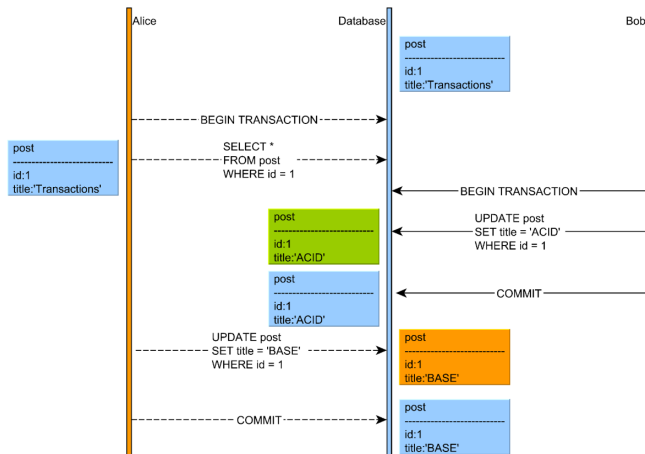


Figure 3: Lost updates

Nonrepeatable reads

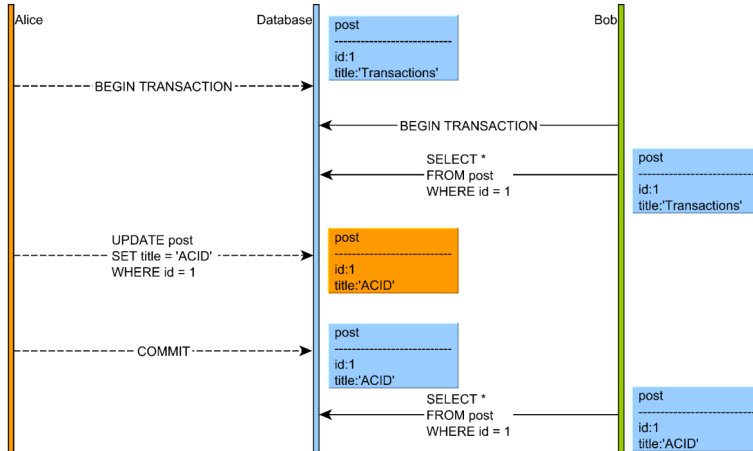


Figure 4: Nonrepeatable reads

Phantom reads

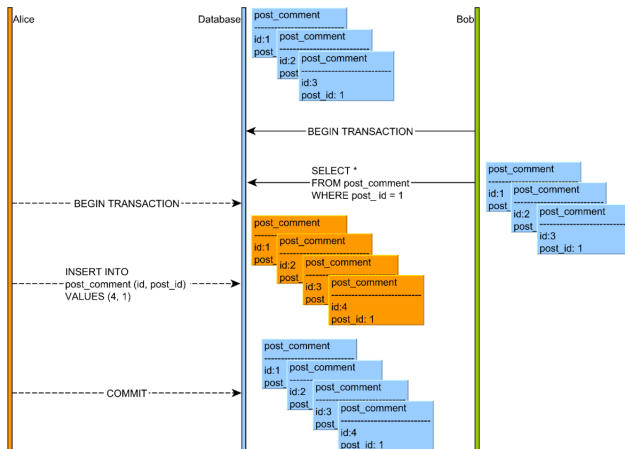


Figure 5: Phantom reads

Concurrency model for disk-based tables

- Pessimistic concurrency control
 - Locking
 - Default for in-the-box SQL server
- Optimistic concurrency control
 - Row-versioning
 - Linked lists of versions
 - Writers use locks, readers don't
 - Default for Azure

Pessimistic concurrency control

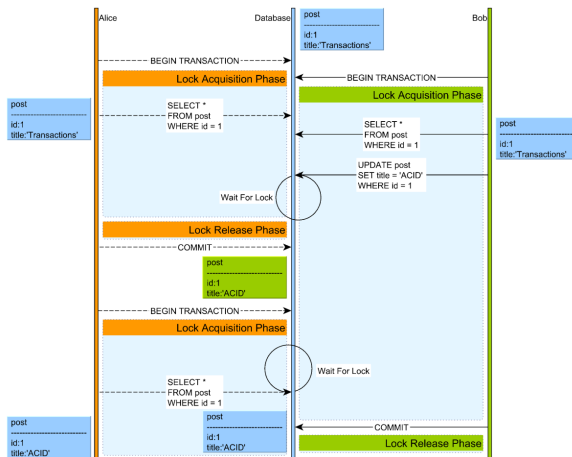


Figure 6: 2PL locking

Optimistic concurrency control

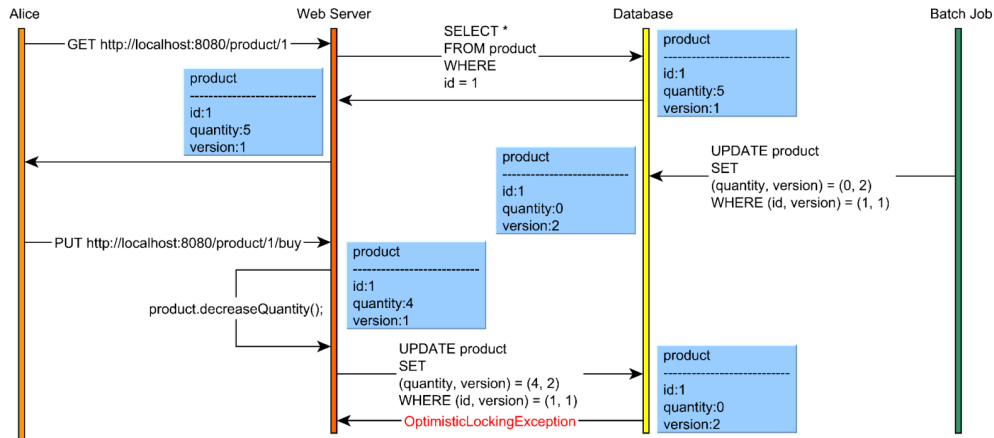


Figure 7: Stateful conversation preventing lost updates

Isolation Level

| Isolation | Dirty reads | Lost updates | Nonrepeatable reads | Phantoms | Concurrency model | Update conflict detection |
|-------------------------|-------------|--------------|---------------------|----------|-------------------|---------------------------|
| Read Uncommitted | Yes | Yes | Yes | Yes | Pessimistic | No |
| Read Committed | No | Yes | Yes | Yes | Pessimistic | No |
| Repeatable Read | No | No | No | Yes | Pessimistic | No |
| Serializable | No | No | No | No | Pessimistic | No |
| Snapshot | No | No | No | No | Optimistic | Yes |
| Read Committed Snapshot | No | Yes | Yes | Yes | Optimistic | No |

Figure 8: Isolation levels

Isolation level (cont.)

- Default values
 - Read committed (in-the-box SQL Server)
 - Read committed snapshot (Azure)
- Set isolation level
 - Use table hint
 - Set at session level
 - `SET ISOLATION LEVEL 'isolation_level'`
 - Row-versioning flags must be set in SQL Server before use

Deadlock

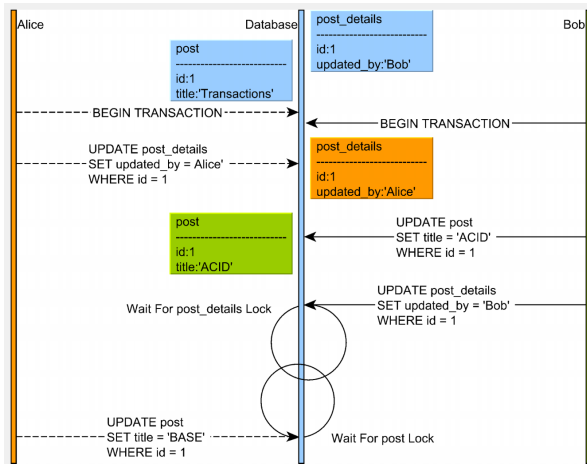


Figure 9: Deadlock

Deadlock prevention coding techniques

- Don't allow transactions to remain open very long
 - keep transactions short
 - keep select statements outside of the transaction except when absolutely necessary
 - never code requests for user input during an open transaction
- Use the lowest possible transaction isolation level
 - default level of read committed is almost always sufficient
 - reserve higher levels for short transactions that make changes to data where integrity is vital
- Make large changes when you can be assured of nearly exclusive access
 - not during peak hours
 - manage to get exclusive access when possible

References

- Bryan Syverson and Murach Joel, SQL Server 2016 for Developers, Mike Murach & Associates Inc., 2016
- Itzik Ben-Gan, Adam Machanic, Dejan Sarka, and Kevin Farlee, T-SQL Querying(Developer Reference), Microsoft Press, 2015
- Vlad Mihalcea, High-Performance Java Persistence, Vlad Mihalcea, 2016